

Philip Greco, III

📍 Tampa, FL ✉ philip.d.greco@gmail.com 🌐 philipgreco.com 📺 philip-greco 🗣 PDGIII

Summary

Senior DevOps Engineer with 10+ years in production infrastructure, including 5+ years designing, automating, and operating large-scale UCaaS platforms across bare-metal Kubernetes and AWS EKS. Drives GitOps-style declarative delivery, infrastructure-as-code, secrets management, and observability, with a track record of leading migrations on live, latency-sensitive VoIP systems.

Experience

Sangoma, Sr. DevOps Engineer

Tampa, FL

July 2025 – present

1 year 1 month

- Own reliability and delivery for a production UCaaS platform spanning bare-metal Kubernetes and AWS EKS clusters, supporting VoIP services including Asterisk, SIP, RTP proxy, and presence.
- Led migration of all manually-managed workloads to a GitOps-style Helm/Helmfile workflow, authoring custom charts and a context guard that eliminated accidental cross-cluster deployments and significantly cut deploy time.
- Led zero-downtime CNI migration to Calico via the Tigera Operator across production clusters; managed Multus and Whereabouts for multi-network VoIP pods.
- Deployed Istio service mesh across clusters, configuring PodDisruptionBudgets, traffic policies, and autoscaling to sustain high availability during rollouts.
- Provisioned AWS EKS clusters with Terraform and Terragrunt, including EFS storage migration, AWS Load Balancer Controller, pod identity for IAM, and VPC security groups.
- Drove adoption of HashiCorp Vault on AWS with External Secrets Operator across all environments, migrating secrets and eliminating hardcoded credentials from Git.
- Built the platform observability stack on EKS with Prometheus, Grafana, Loki, and Alloy using daemonset and sidecar logging patterns, reducing mean time to detection.
- Automated Kubernetes upgrades across all clusters through v1.35 via Ansible playbooks with pod health checks, DNS validation, and controlled rolling upgrades, with zero customer-facing downtime.
- Serve in the on-call rotation, leading incident response and root cause analysis for production escalations and authoring runbooks adopted by the team.

Sangoma, DevOps Engineer

Atlanta, GA

Jan 2021 – July 2025

4 years 7 months

- Provisioned bare-metal Dell blade servers via PXE boot with Debian preseed automation, handling network configuration and OS deployment end to end.
- Authored Ansible roles to compile and deploy FreeSWITCH from source across the fleet, managing version pinning, Sofia-SIP patches, and debug symbol builds.
- Managed fleets of SIP and global load balancers via Ansible, maintaining version consistency and high availability in production.
- Performed root cause analysis on production issues escalated by Tier 3 Support, resolving escalations and feeding fixes back into automation.
- Built internal Bash and SQL diagnostic tooling adopted across the support organization, reducing mean time to resolution.

Star2Star Communications, Lead Tier 3 Technician

Atlanta, GA

Jan 2016 – Jan 2021

5 years 1 month

- Diagnosed and resolved complex SIP/VoIP issues using Wireshark packet analysis, filing detailed bug reports with captures and reproduction steps.
- Correlated customer logs and packet captures to identify platform-wide SIP issues, escalating to development and DevOps teams.
- Authored technical documentation and Standard Operating Procedures adopted across the support organization.
- Mentored junior technicians on SIP troubleshooting and diagnostic tooling.

Projects

Self-Hosted Kubernetes Homelab

- Designed and operate a bare-metal Kubernetes cluster managed entirely through GitOps with Helmfile — declarative, dependency-ordered releases with Git as the single source of truth.
- Integrated HashiCorp Vault with External Secrets Operator via Kubernetes auth for dynamic secret injection without storing credentials in Git.
- Deployed the Tailscale operator with custom ingress for zero-config HTTPS access to internal services without exposing public ports.
- Built an observability stack with Prometheus, Grafana, Loki, Alloy, and Alertmanager for metrics, logs, and alerting with SLO-based alert routing.
- Provisioned Longhorn distributed block storage for stateful workloads and MetalLB for bare-metal load balancing.
- Automated CI/CD with GitHub Actions and self-hosted ARC runner scale sets, publishing images to GHCR and deploying via Helm.

Self-Hosted LLM Inference & AI Automation

- Self-hosted Ollama on Kubernetes for fully local LLM inference, keeping all data on-premises with no external AI APIs.
- Built a Python RAG pipeline using ChromaDB and Ollama-generated embeddings for semantic search over historical articles, injecting context into prompts.
- Deployed as a Kubernetes CronJob with scheduling jitter, concurrency controls, and credentials pulled from Vault via External Secrets Operator.

Education

BS **Southern New Hampshire University**, Information Technology
• Concentration: Networking & Telecommunications

Manchester, NH

Jan 2013 – Jan 2017

AS **NHTI, Concord's Community College**, Information Technology

Concord, NH

Jan 2011 – Jan 2013

Skills

Container & Orchestration: Kubernetes, Docker, EKS, Helm, Helmfile

Infrastructure as Code & GitOps: Terraform, Terragrunt, Ansible, Helmfile, declarative Git-driven delivery

CI/CD: GitHub Actions, Actions Runner Controller (ARC), GHCR

Secrets & Security: HashiCorp Vault, External Secrets Operator, Kubernetes RBAC, Tailscale

Observability & Reliability: Prometheus, Grafana, Loki, Alloy, Alertmanager, incident response, root cause analysis

Networking & Ingress: Istio, Traefik, MetalLB, NGINX, Wireguard, VoIP/SIP, TCP/IP, DNS

Storage: Longhorn, AWS EFS, S3

Cloud & OS: AWS, Digital Ocean, Linux (Ubuntu, Debian)

Languages: Python, Bash, SQL

AI & LLM: Ollama, ChromaDB, RAG pipelines